



LECTURE 4 MACHINE LEARNINIG

Win+w



Outline




- Introduction
- Course Outline
- Example Implementation
- Summary




Example Applications of Graph Analytics

ISR




- Graphs represent entities and relationships detected through multi-INT sources
- 1,000s – 1,000,000s tracks and locations
- GOAL: Identify anomalous patterns of life

Social



- Graphs represent relationships between individuals or documents
- 10,000s – 10,000,000s individual and interactions
- GOAL: Identify hidden social networks

Cyber




- Graphs represent communication patterns of computers on a network
- 1,000,000s – 1,000,000,000s network events
- GOAL: Detect cyber attacks or malicious software

- Cross-Mission Challenge: Detection of subtle patterns in massive multi-source noisy datasets



Example: Web Traffic Graph



Graph Statistics

- 90 minutes worth of traffic
- 1 frame = 1 minute of traffic
- Number of source computers: 4,063
- Number of web servers: 16,397
- Number of logs: 4,344,148

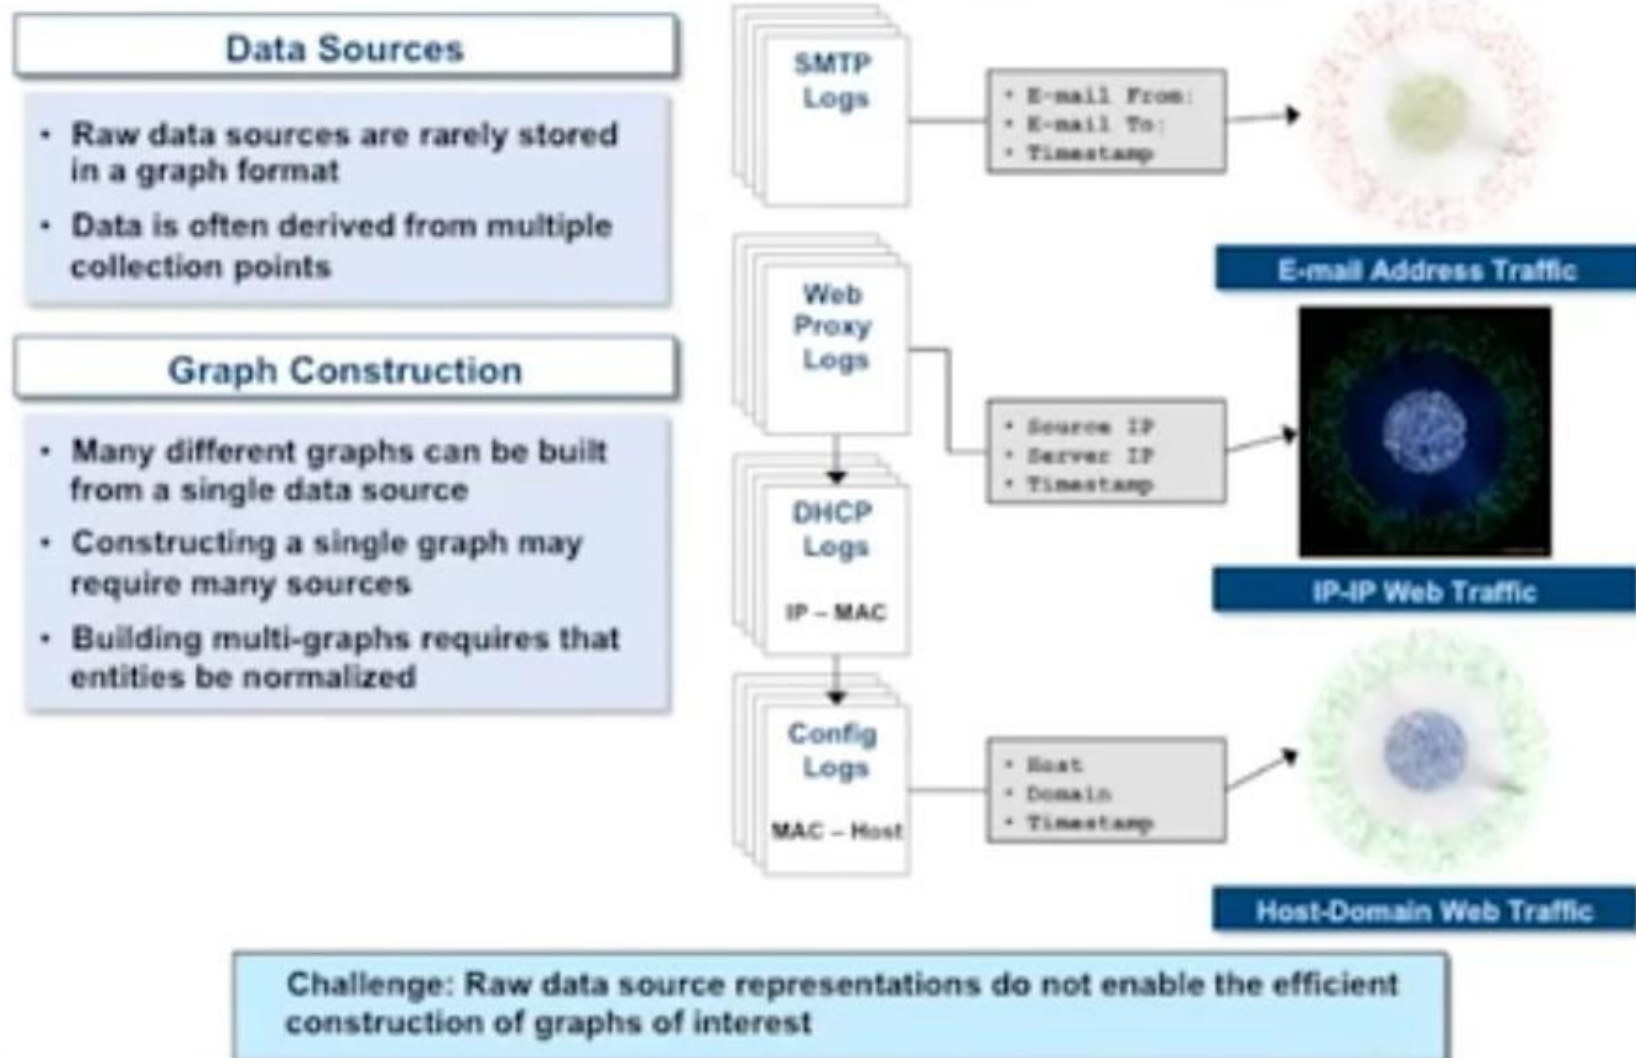
Malicious Activity Statistics

- Number of infected IPs: 1
- Number of event logs: 16,000
- % infected traffic: 0.37%
- Existing tools did not detect event
- Detection took **10 days** and required manual log inspection

Challenge: Activity signature is typically a weak signal

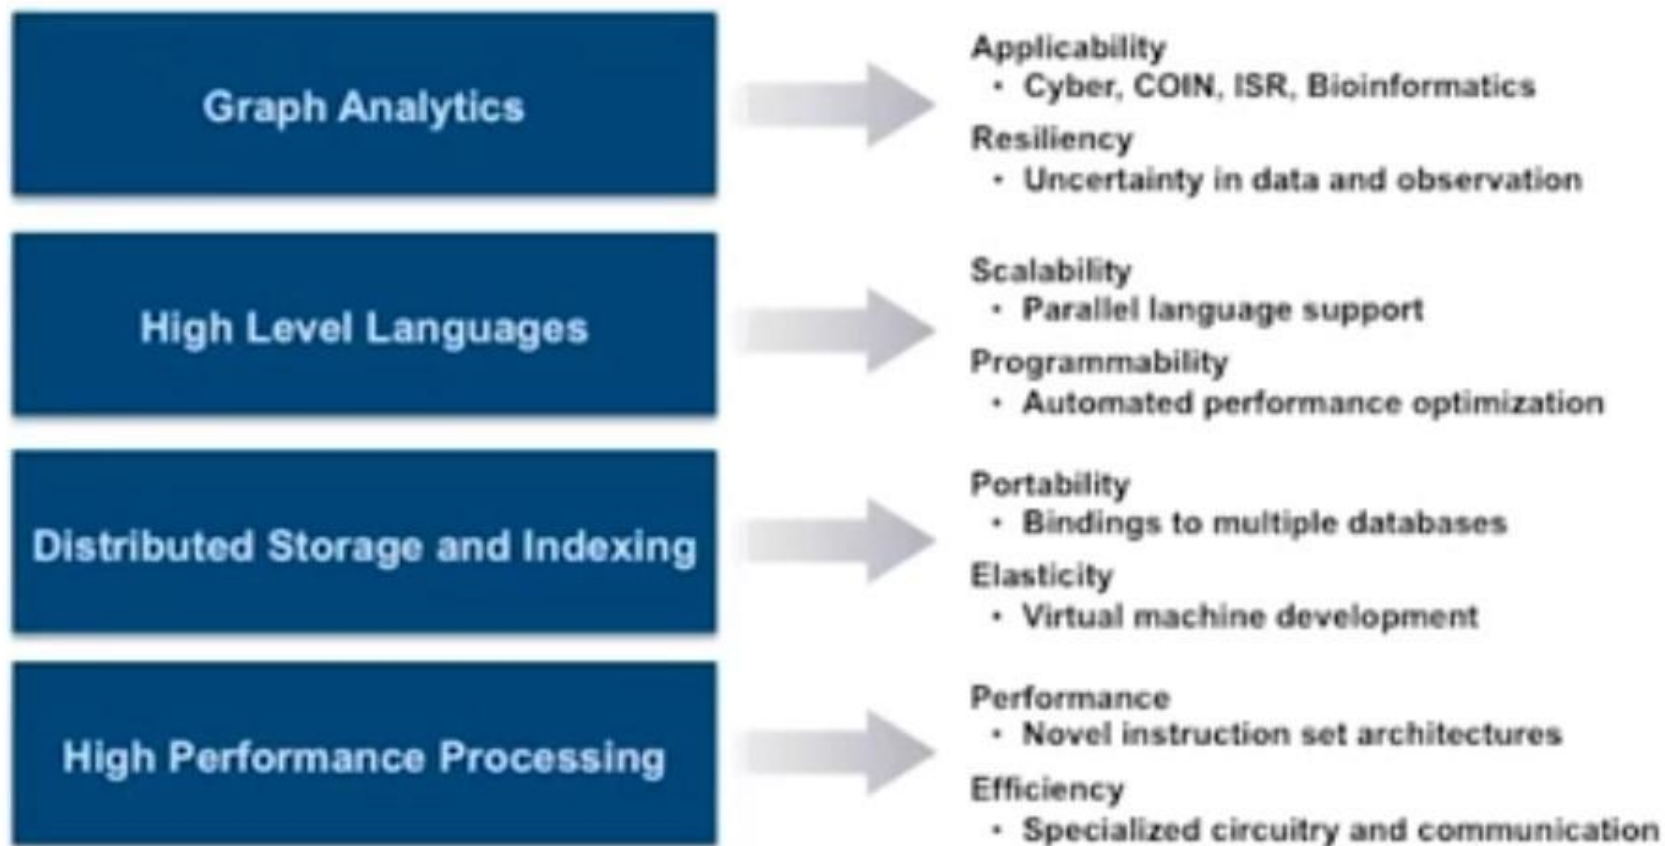


Big Data Challenge: Data Representation



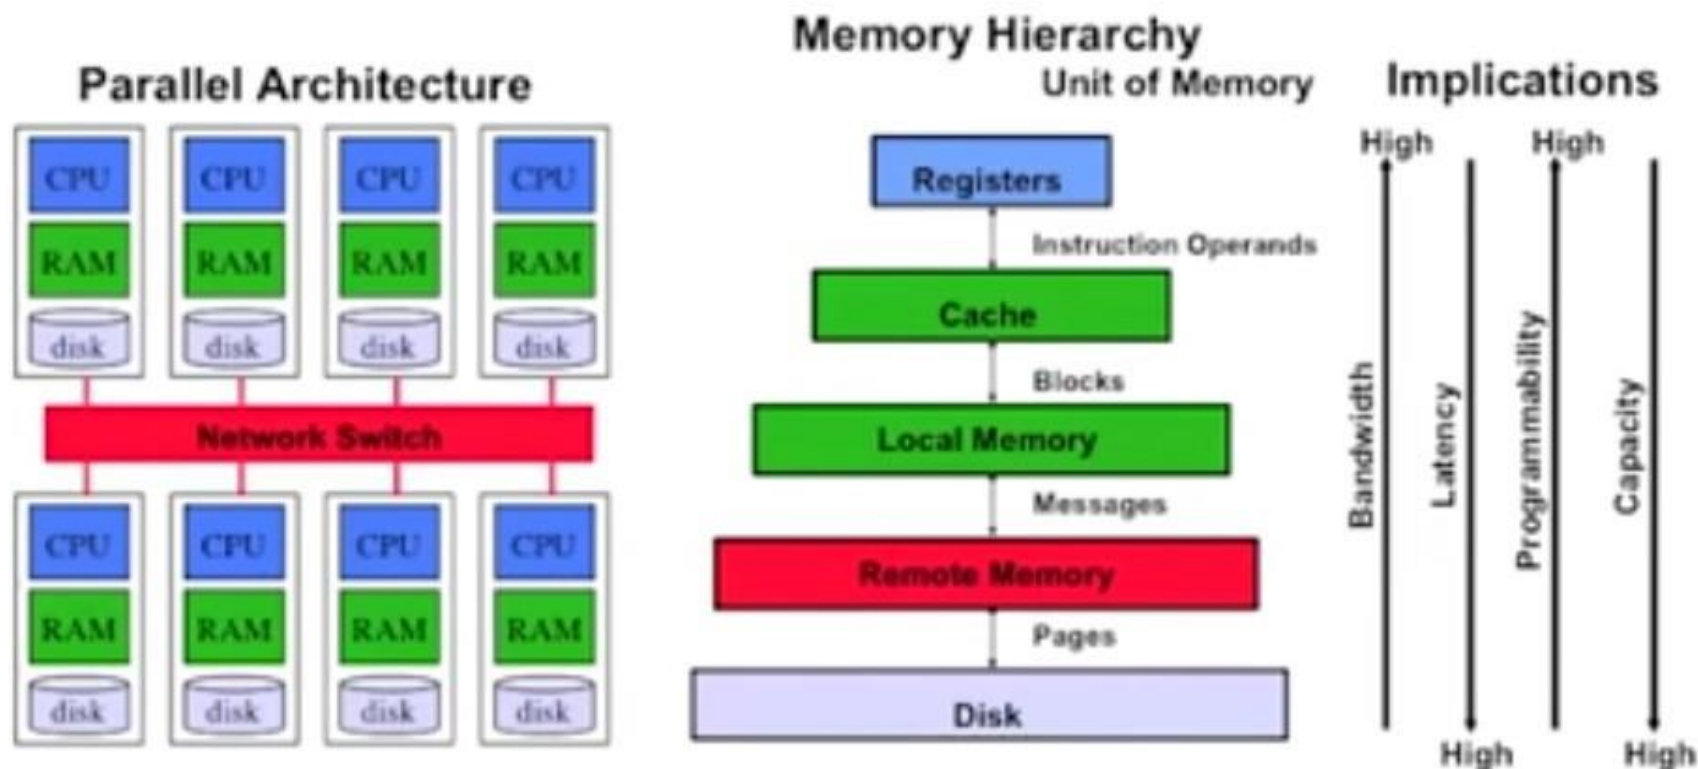


Technology Stack





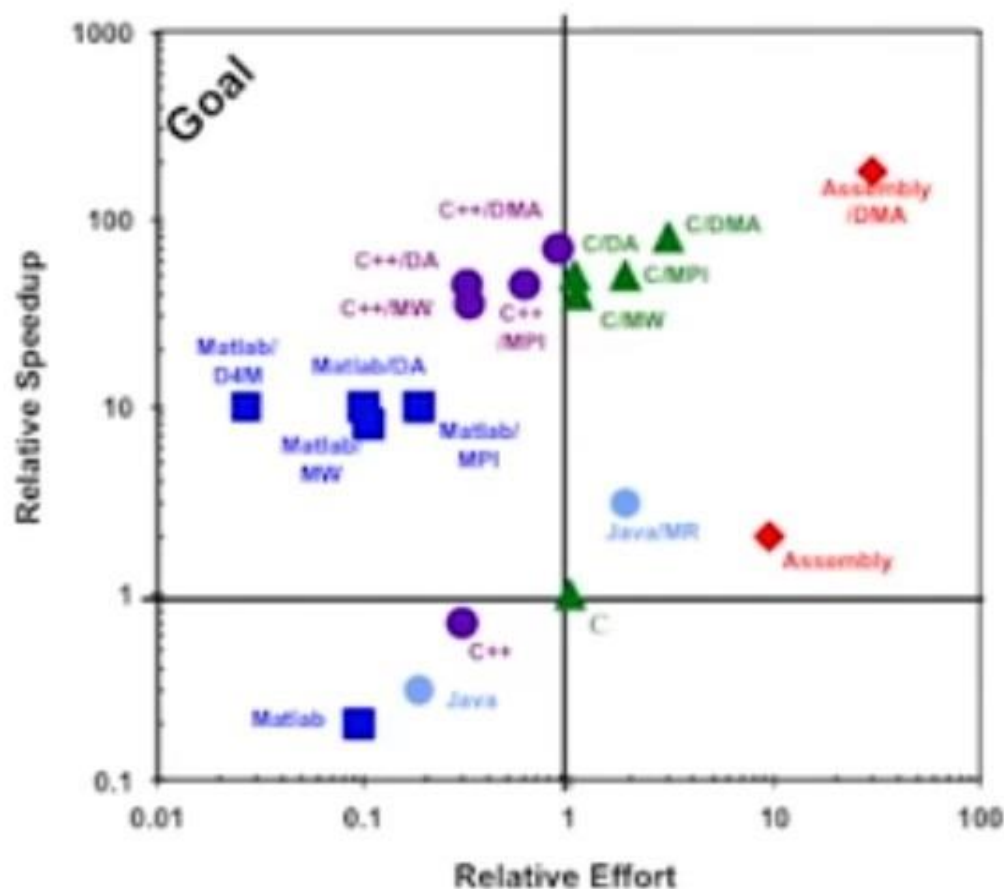
Software and Bytes Live on Parallel Computers



- Nearly all modern computers are Von Neumann architectures with multi-level memory hierarchies
- The architecture selects the algorithms and data that run well on it



Software Performance vs. Parallel Programmer Effort




Parallel Programming Models

| | |
|-----|--|
| DMA | = Direct Memory Access |
| MPI | = Message Passing |
| DA | = Distributed Arrays |
| MW | = Manager/Worker |
| MR | = Map/Reduce |
| D4M | = Dynamic Distributed Dimensional Data Model |

- Goal: Software that does a lot with the least effort




Data Use Cases



- Data volume and data request size determine best approach
- Always want to start with the simplest and move to the most complex



The Fast Path



- The class teaches the highest performance and lowest effort software techniques that are currently known



Key Course Concepts

- Bigger definition of a graph
 - How to move beyond random, undirected, unweighted graphs to power-law, directed, multi-hyper graphs
- Bigger definition of linear algebra
 - How to move beyond real numbers to doing math with words and strings
- Bigger definition of processing
 - How to move beyond map/reduce to distributed arrays programming

• These abstract concepts are the foundation for high performance signal processing on large unstructured data sets




Course Outline

- **Introduction**
 - Review course goals and structure
- **Using Associative Arrays**
 - Schemas, incidence matrices, and directed multi-hyper graphs
- **Group Theory**
 - Extending linear algebra to words using fuzzy algebra
- **Entity Analysis in Unstructured Data**
 - Reading and parsing unstructured data
- **Analysis of Structured Data**
 - Graph traversal queries
- **Power Law Data**
 - Models and fitting
- **Cross Correlation**
 - Sequence data, computing degree distributions, and finding matches
- **Parallel Processing**
 - Kronecker graphs, parallel data generation and computation
- **Databases**
 - Relational, triple store, and exploded schemas



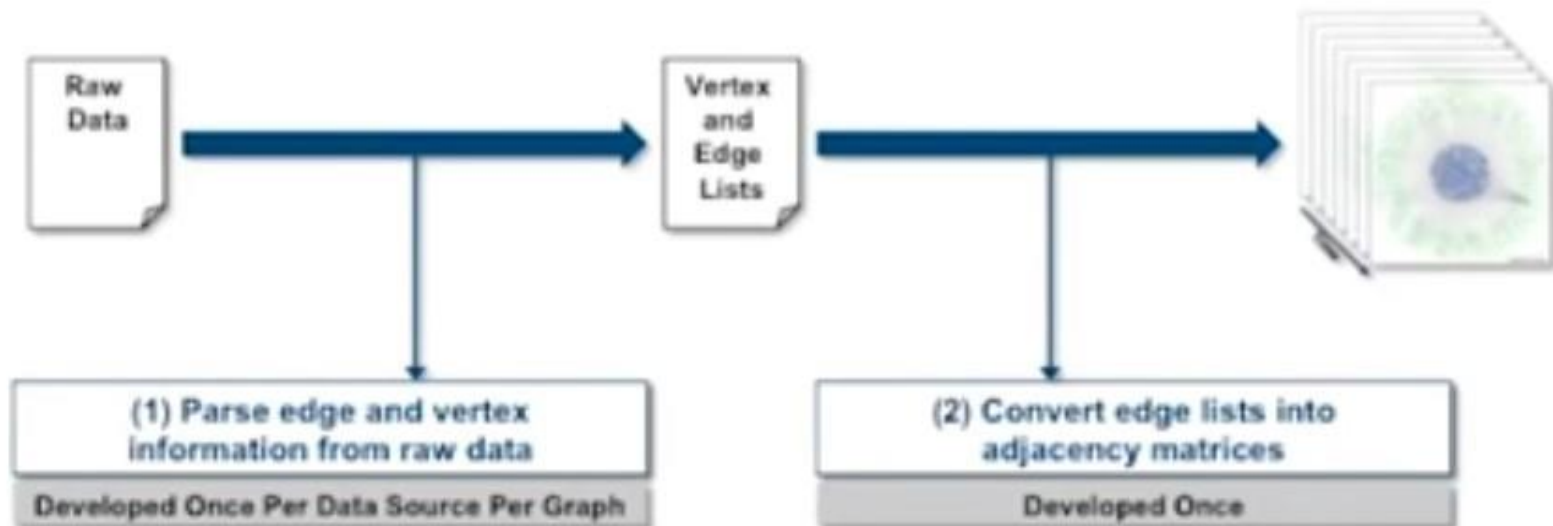
References

- Book: "Graph Algorithms in the Language of Linear Algebra"
- Editors: Kepner (MIT-LL) and Gilbert (UCSB)
- Contributors:
 - Bader (Ga Tech)
 - Bliss (MIT-LL)
 - Bond (MIT-LL)
 - Dunlavy (Sandia)
 - Faloutsos (CMU)
 - Fineman (CMU)
 - Gilbert (USCB)
 - Heitsch (Ga Tech)
 - Hendrickson (Sandia)
 - Kegelmeyer (Sandia)
 - Kepner (MIT-LL)
 - Kolda (Sandia)
 - Leskovec (CMU)
 - Madduri (Ga Tech)
 - Mohindra (MIT-LL)
 - Nguyen (MIT)
 - Radar (MIT-LL)
 - Reinhardt (Microsoft)
 - Robinson (MIT-LL)
 - Shah (USCB)





Constructing Graph Representations of Raw Data Source

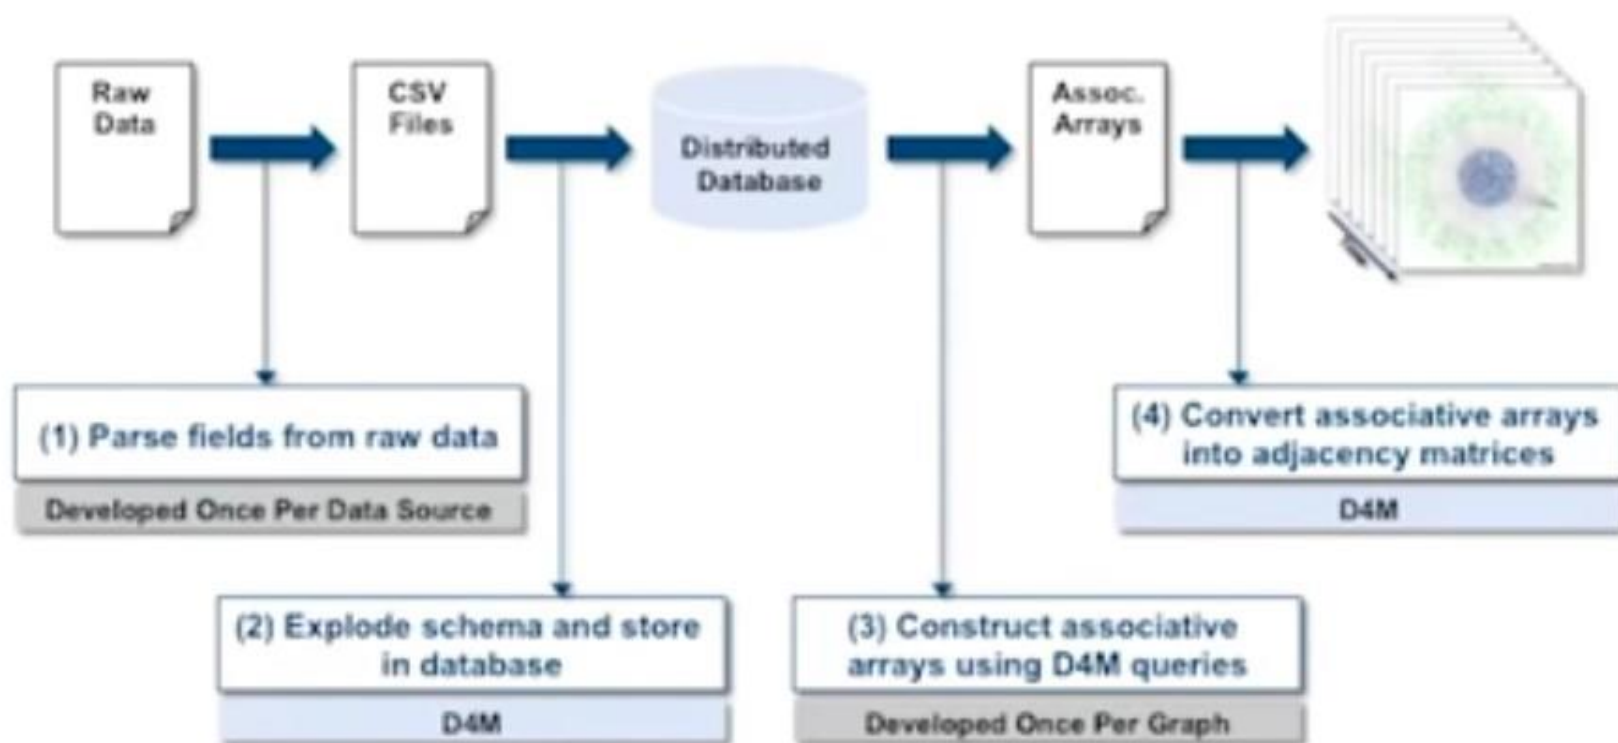


- Raw data sources can contain information about multiple types of relations between entities
- The process of constructing a graph representation is specific to both the data source and the relationships represented by the graph

- The development time of parsing and graph construction algorithms can overwhelm the runtime of the algorithm




Graph Construction Using D4M



- D4M provides needed flexibility in the construction of large-scale, dynamic graphs at different resolutions and scopes



Graph Construction Using D4M: Parsing Raw Data Into Dense Tables



Proxy Logs


```
128.0.0.1 208.29.69.138 "-" [10/May/2011:09:52:53] "GET http://www.thedailybeast.com/ HTTP/1.1" 200  
1024 6192 "http://www.theatlantic.com/" "Mozilla/5.0 (X11; U; Linux x86_64; en-US; rv:1.9.2.13)  
Gecko/20101209 CentOS/3.6-2.el5.centos Firefox/3.6.13" "b1" - "text/html" "MITLAB" 0.523 "-"  
Neutral TCP_MISS  
192.168.1.1 157.166.255.18 "-" [12/May/2011:13:24:11] "GET http://www.cnn.com/ HTTP/1.1" 335 256  
10296 "-" "Mozilla/5.0 (X11; U; Linux x86_64; en-US; rv:1.9.2.13) Gecko/20101209 CentOS/  
3.6-2.el5.centos Firefox/3.6.13" "bu" - "text/html" "MITLAB" 0.784 "-" Neutral TCP_MISS  
...
```

Dense Table

| log_id | src_ip | server_ip | time_stamp | req_line | ... |
|--------|-------------|----------------|----------------------|--|-----|
| 001 | 128.0.0.1 | 208.29.69.138 | 10/May/2011:09:52:53 | GET http://www.thedailybeast.com/ HTTP/1.1 | ... |
| 002 | 192.168.1.2 | 157.166.255.18 | 12/May/2011:13:24:11 | GET http://www.cnn.com/ HTTP/1.1 | ... |
| 003 | 128.0.0.1 | 74.125.224.72 | 13/May/2011:11:05:12 | GET http://www.google.com/ HTTP/1.1 | ... |
| ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ |



Graph Construction Using D4M: Explode Schema





Graph Construction Using D4M: Storing Exploded Data as Triples



Exploded Table

| | src_ip[128.0.0.1] | src_ip[192.168.1.2] | server_ip[157.166.255.18] | server_ip[208.29.69.138] | server_ip[74.125.224.72] |
|-------------|-------------------|---------------------|---------------------------|--------------------------|--------------------------|
| log_id[001] | 1 | 0 | 0 | 1 | 0 |
| log_id[002] | 0 | 1 | 1 | 0 | 0 |
| log_id[003] | 1 | 0 | 0 | 0 | 1 |

D4M stores the triple data representing both
the exploded table and its transpose

Table Triples

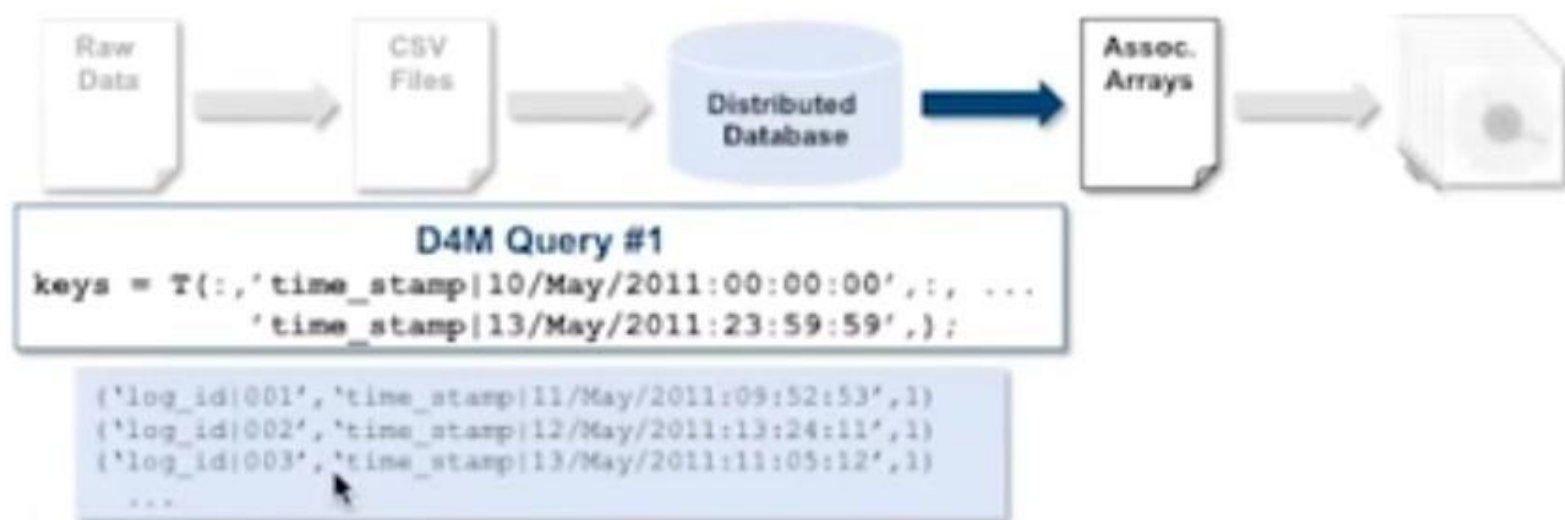
| Row | Column | Value |
|-------------|---------------------------|-------|
| log_id[001] | src_ip[128.0.0.1] | 1 |
| log_id[001] | server_ip[208.29.69.138] | 1 |
| log_id[002] | src_ip[192.168.1.2] | 1 |
| log_id[002] | server_ip[157.166.255.18] | 1 |
| log_id[003] | src_ip[128.0.0.1] | 1 |
| log_id[003] | server_ip[74.125.224.72] | 1 |

Table Transpose Triples

| Row | Column | Value |
|---------------------------|-------------|-------|
| server_ip[157.166.255.18] | log_id[002] | 1 |
| server_ip[208.29.69.138] | log_id[001] | 1 |
| server_ip[74.125.224.72] | log_id[003] | 1 |
| src_ip[128.0.0.1] | log_id[001] | 1 |
| src_ip[128.0.0.1] | log_id[003] | 1 |
| src_ip[192.168.1.2] | log_id[002] | 1 |




Graph Construction Using D4M: Construct Associative Arrays






Graph Construction Using D4M: Construct Associative Arrays



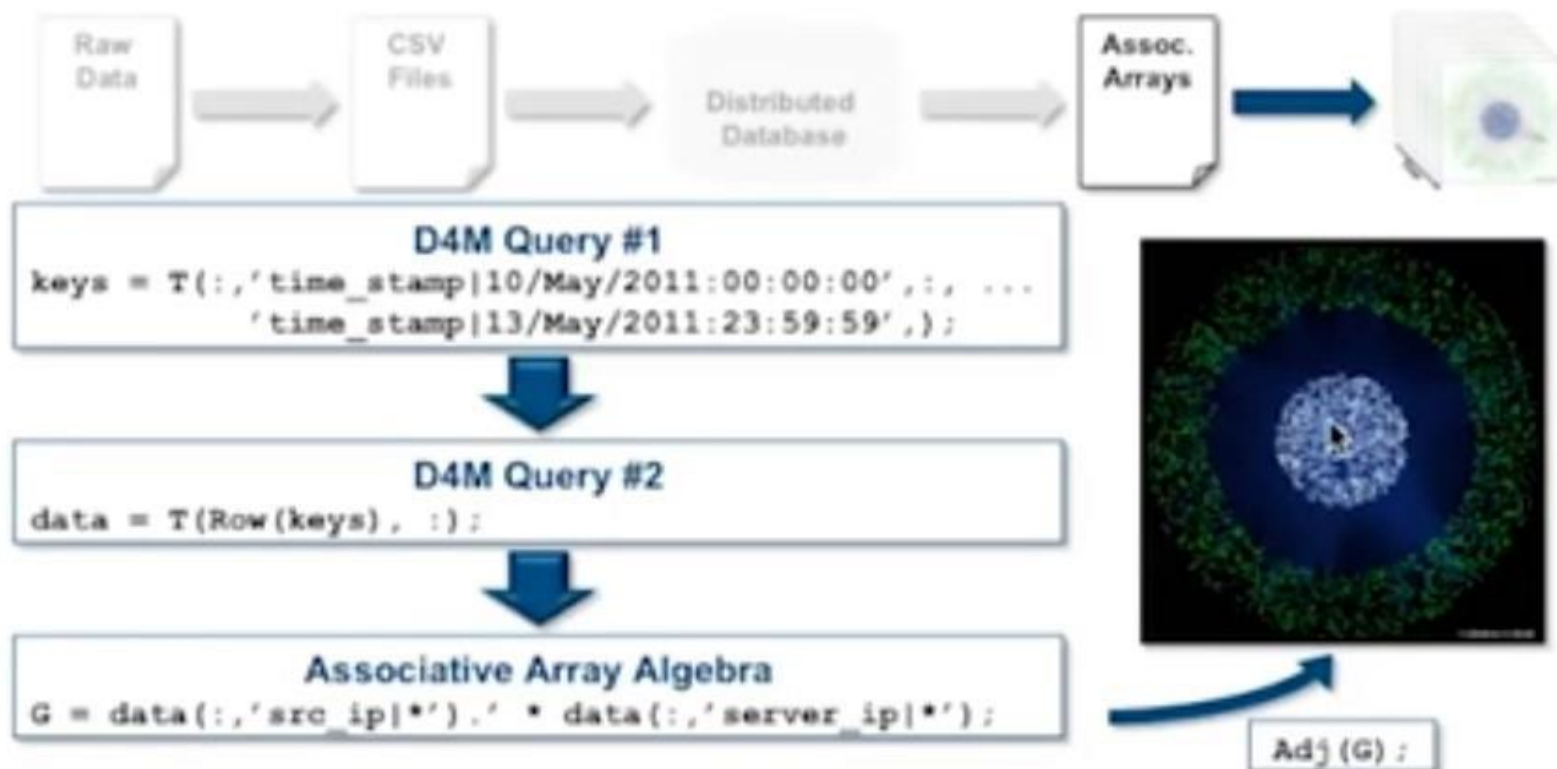


Graph Construction Using D4M: Construct Associative Arrays






Graph Construction Using D4M: Construct Associative Arrays



- Graphs can be constructed with minimal effort using D4M queries and associative array algebra



Constructing Graph Representation of One Week's Worth of Proxy Data



- Ingested ~130 million proxy log records resulting in ~4.5 billion triples
- Constructed 604,800 secondwise source IP to server IP graphs
- Constructing graphs with different vertex types could be done without re-parsing or re-ingesting data

- Utilizing D4M could allow analysis to be run in nearly real-time (dependent on raw data availability)



Summary

- Big data is found across a wide range of areas
 - Document analysis
 - Computer network analysis
 - DNA Sequencing
- Currently there is a gap in big data analysis tools for algorithm developers
- D4M fills this gap by providing algorithm developers composable associative arrays that admit linear algebraic manipulation



Outline




- Introduction
- Course Outline
- Example Implementation
- Summary




Example Applications of Graph Analytics

ISR




- Graphs represent entities and relationships detected through multi-INT sources
- 1,000s – 1,000,000s tracks and locations
- GOAL: Identify anomalous patterns of life

Social



- Graphs represent relationships between individuals or documents
- 10,000s – 10,000,000s individual and interactions
- GOAL: Identify hidden social networks

Cyber




- Graphs represent communication patterns of computers on a network
- 1,000,000s – 1,000,000,000s network events
- GOAL: Detect cyber attacks or malicious software

- Cross-Mission Challenge: Detection of subtle patterns in massive multi-source noisy datasets



Example: Web Traffic Graph



Graph Statistics

- 90 minutes worth of traffic
- 1 frame = 1 minute of traffic
- Number of source computers: 4,063
- Number of web servers: 16,397
- Number of logs: 4,344,148

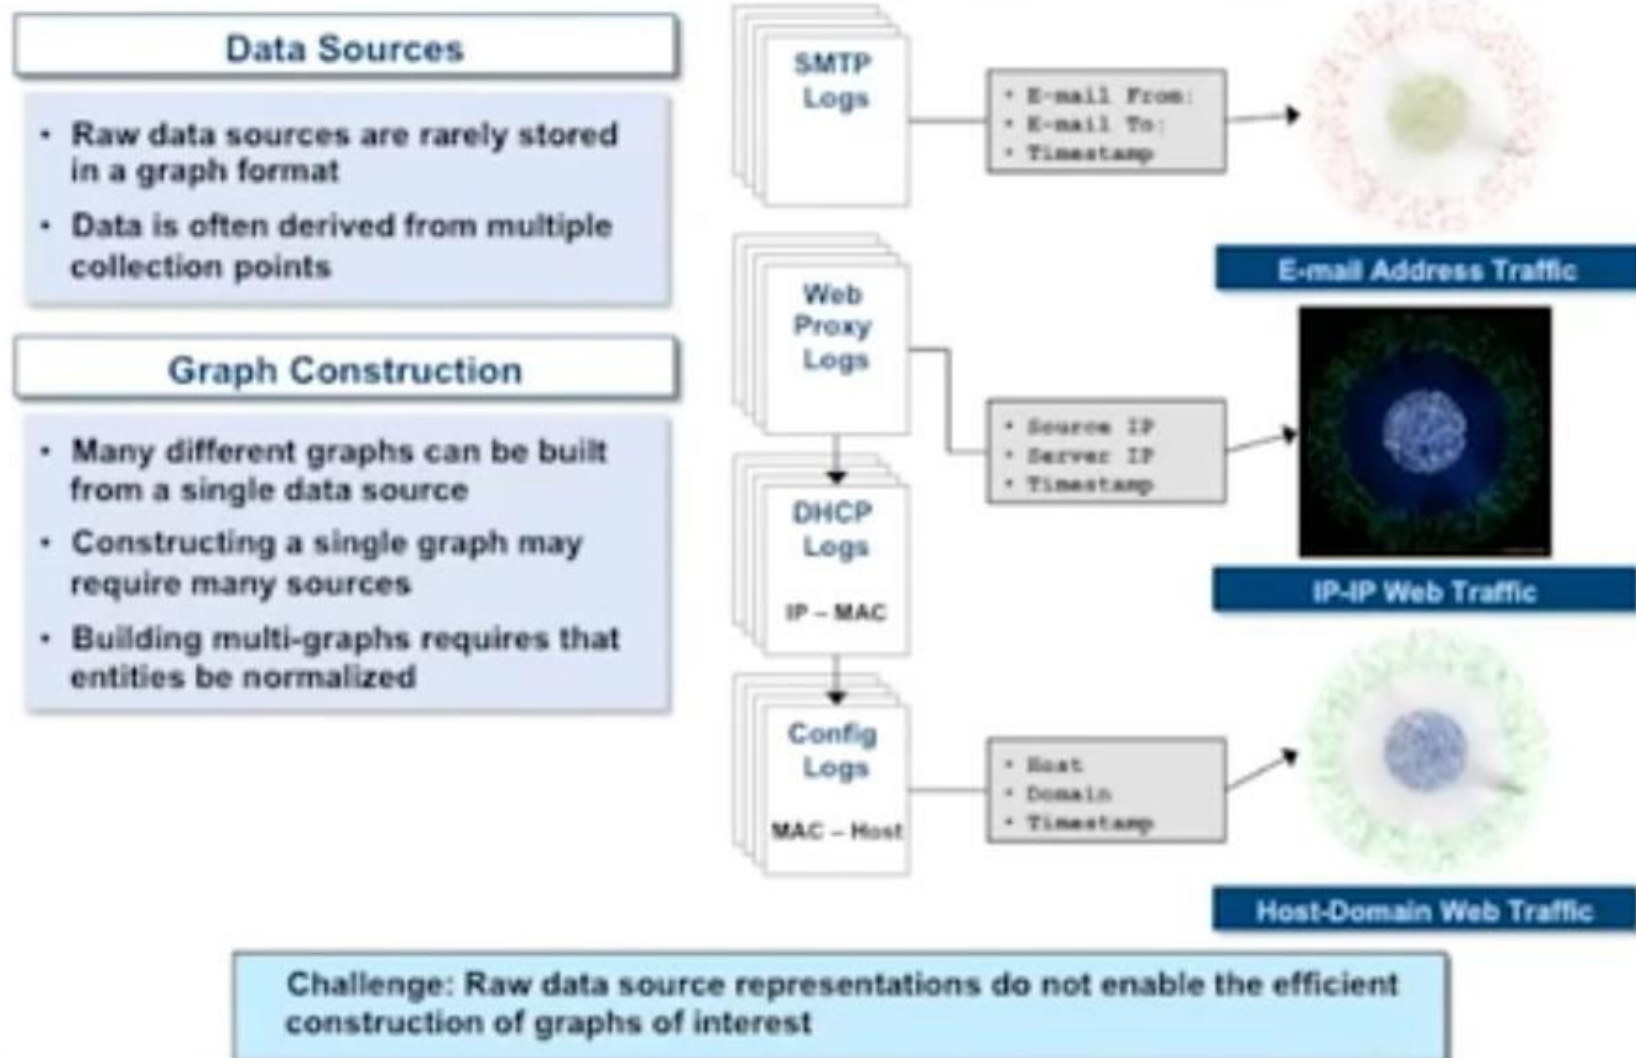
Malicious Activity Statistics

- Number of infected IPs: 1
- Number of event logs: 16,000
- % infected traffic: 0.37%
- Existing tools did not detect event
- Detection took **10 days** and required manual log inspection

Challenge: Activity signature is typically a weak signal

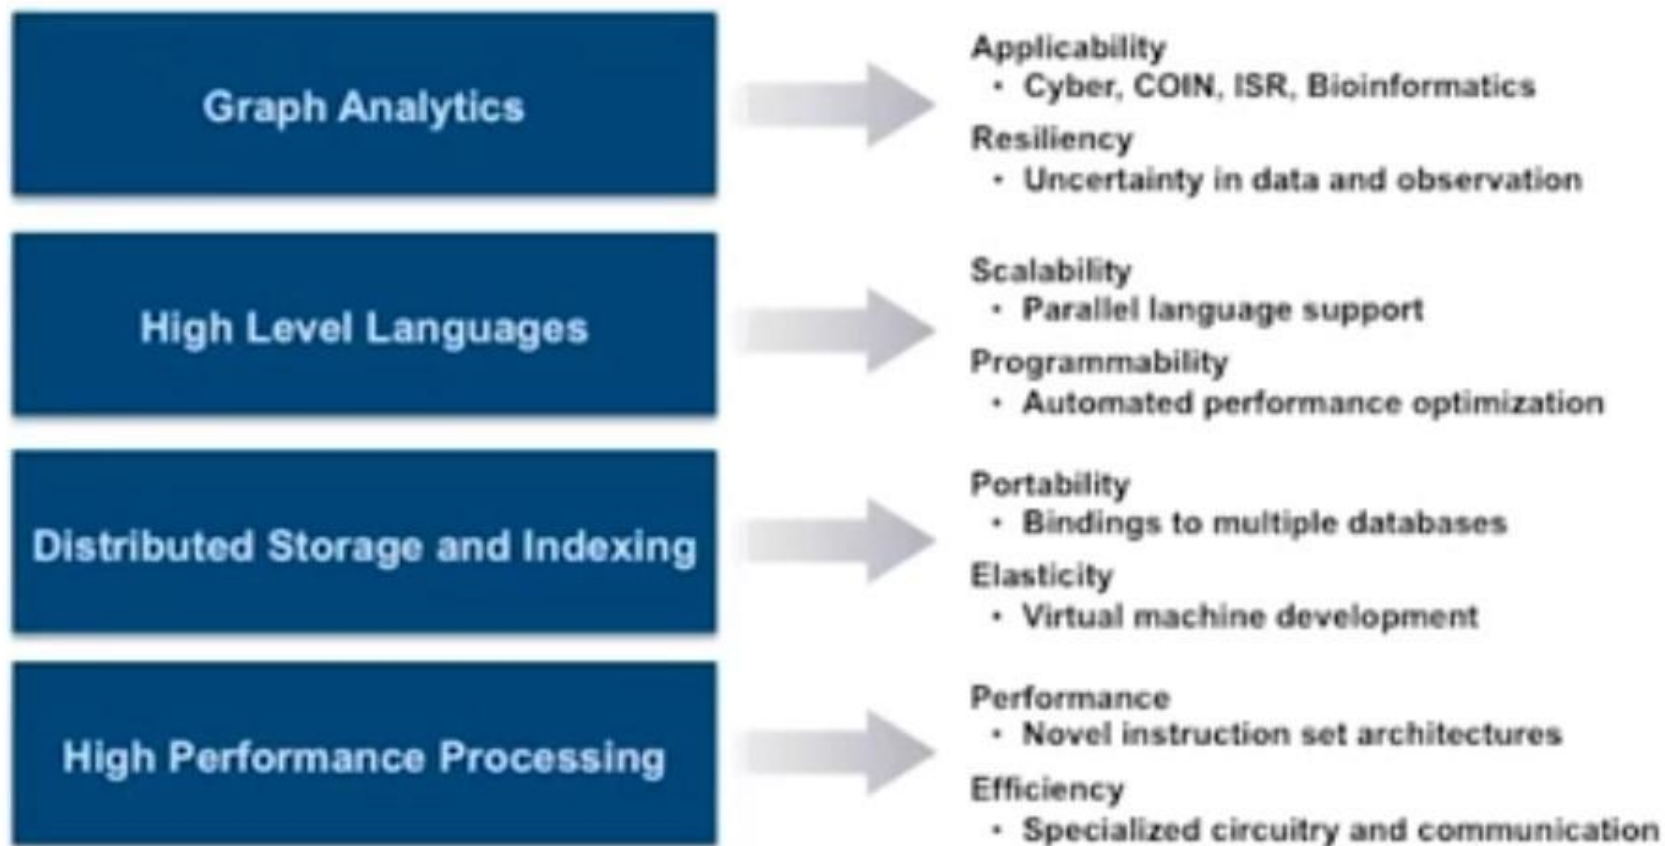


Big Data Challenge: Data Representation



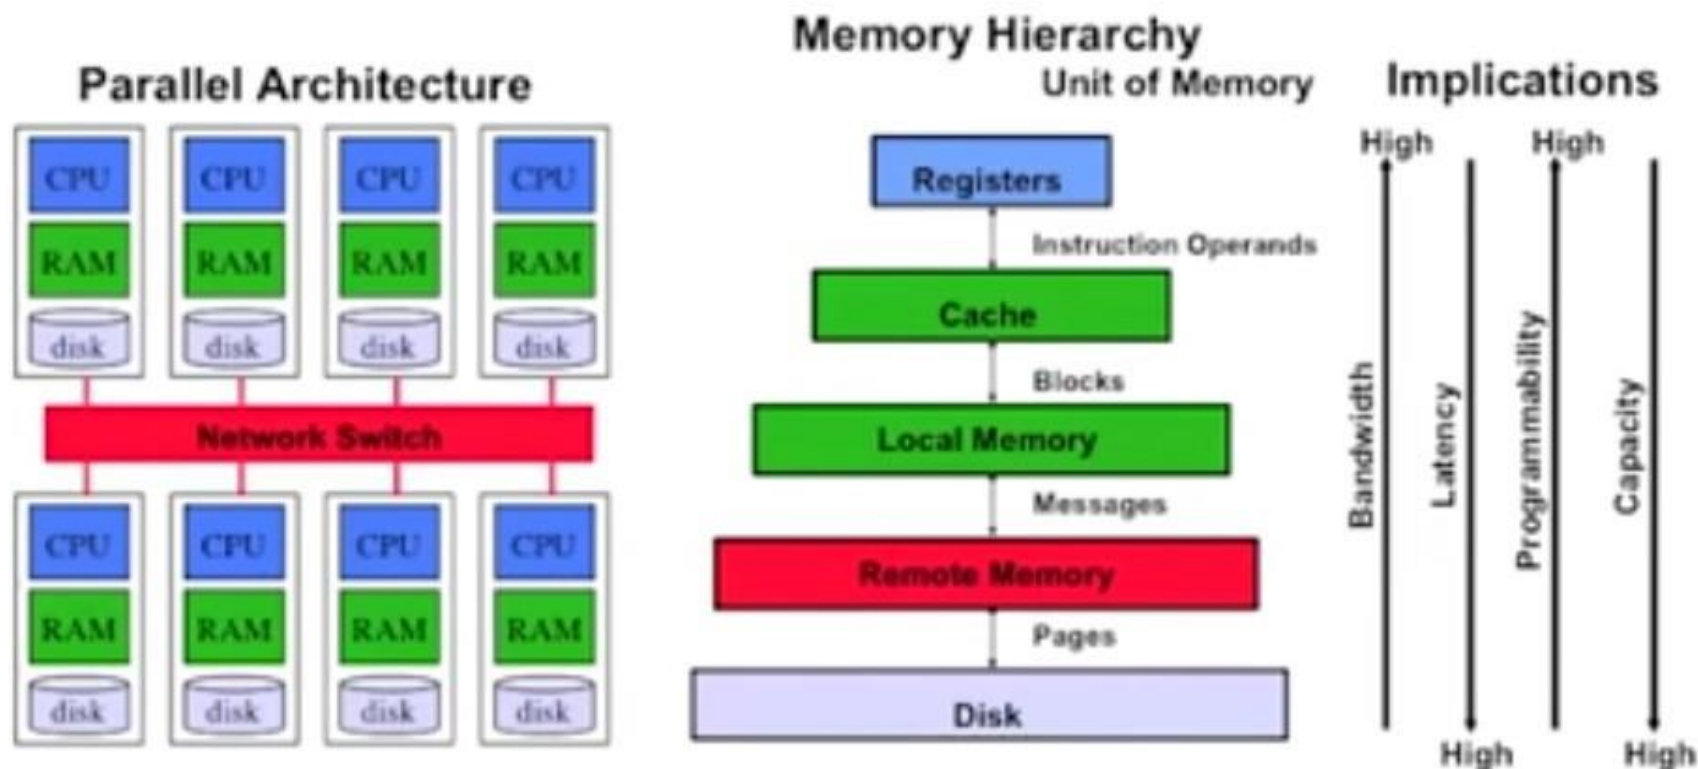


Technology Stack





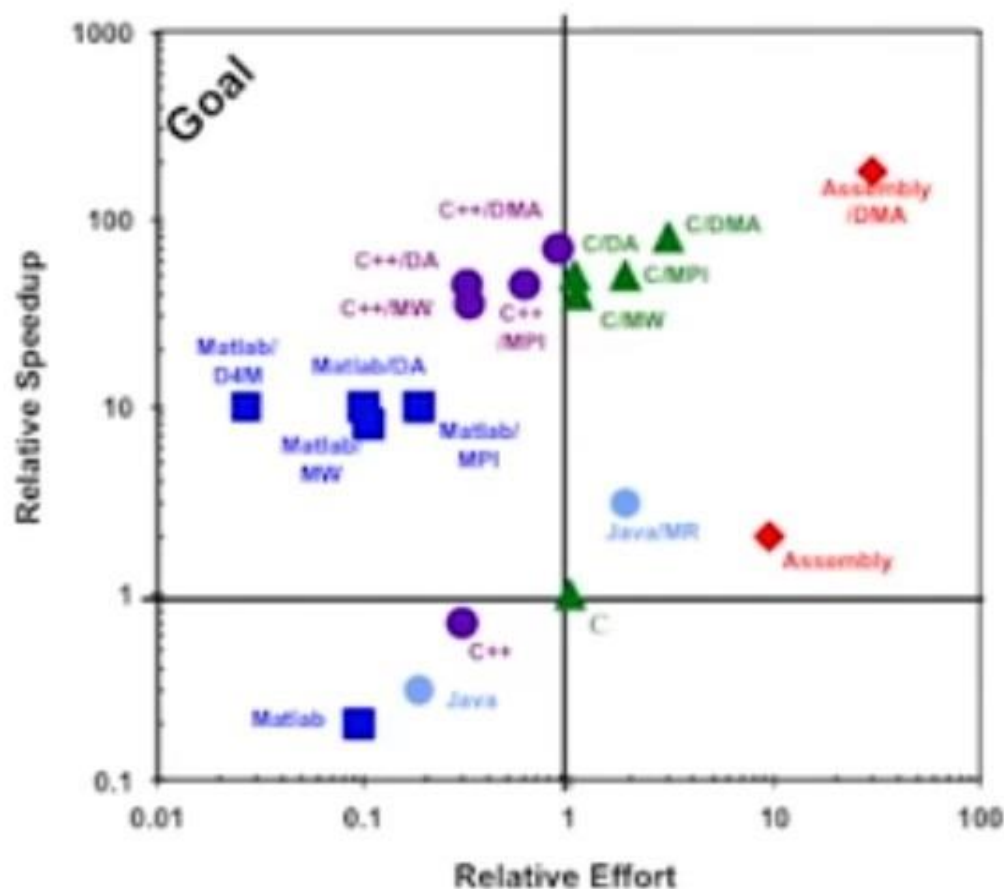
Software and Bytes Live on Parallel Computers



- Nearly all modern computers are Von Neumann architectures with multi-level memory hierarchies
- The architecture selects the algorithms and data that run well on it



Software Performance vs. Parallel Programmer Effort



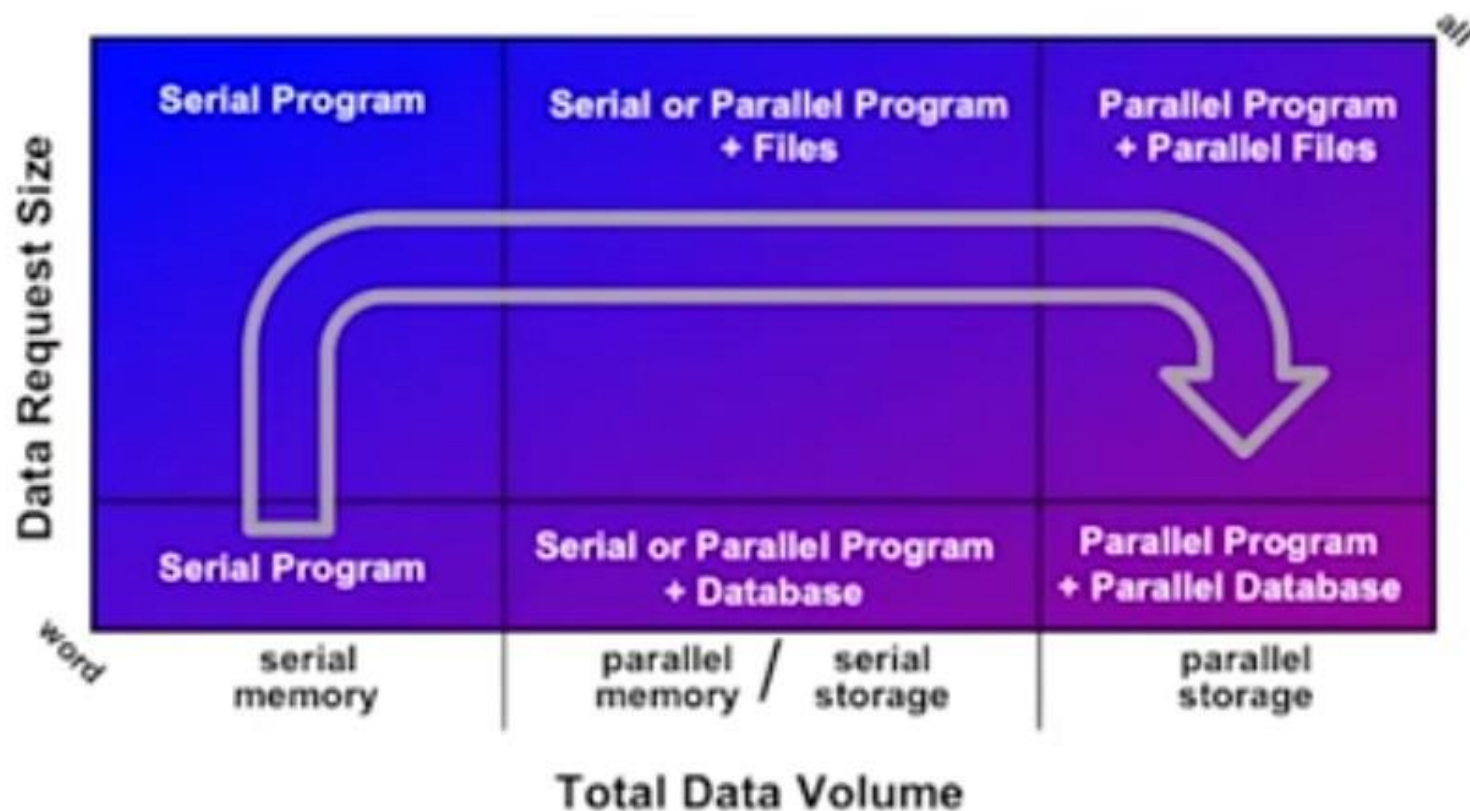
Parallel Programming Models

| | |
|-----|--|
| DMA | = Direct Memory Access |
| MPI | = Message Passing |
| DA | = Distributed Arrays |
| MW | = Manager/Worker |
| MR | = Map/Reduce |
| D4M | = Dynamic Distributed Dimensional Data Model |

- Goal: Software that does a lot with the least effort




Data Use Cases



- Data volume and data request size determine best approach
- Always want to start with the simplest and move to the most complex



The Fast Path



- The class teaches the highest performance and lowest effort software techniques that are currently known



Key Course Concepts

- Bigger definition of a graph
 - How to move beyond random, undirected, unweighted graphs to power-law, directed, multi-hyper graphs
- Bigger definition of linear algebra
 - How to move beyond real numbers to doing math with words and strings
- Bigger definition of processing
 - How to move beyond map/reduce to distributed arrays programming

• These abstract concepts are the foundation for high performance signal processing on large unstructured data sets




Course Outline

- **Introduction**
 - Review course goals and structure
- **Using Associative Arrays**
 - Schemas, incidence matrices, and directed multi-hyper graphs
- **Group Theory**
 - Extending linear algebra to words using fuzzy algebra
- **Entity Analysis in Unstructured Data**
 - Reading and parsing unstructured data
- **Analysis of Structured Data**
 - Graph traversal queries
- **Power Law Data**
 - Models and fitting
- **Cross Correlation**
 - Sequence data, computing degree distributions, and finding matches
- **Parallel Processing**
 - Kronecker graphs, parallel data generation and computation
- **Databases**
 - Relational, triple store, and exploded schemas




References

- Book: "Graph Algorithms in the Language of Linear Algebra"
- Editors: Kepner (MIT-LL) and Gilbert (UCSB)
- Contributors:
 - Bader (Ga Tech)
 - Bliss (MIT-LL)
 - Bond (MIT-LL)
 - Dunlavy (Sandia)
 - Faloutsos (CMU)
 - Fineman (CMU)
 - Gilbert (USCB)
 - Heitsch (Ga Tech)
 - Hendrickson (Sandia)
 - Kegelmeyer (Sandia)
 - Kepner (MIT-LL)
 - Kolda (Sandia)
 - Leskovec (CMU)
 - Madduri (Ga Tech)
 - Mohindra (MIT-LL)
 - Nguyen (MIT)
 - Radar (MIT-LL)
 - Reinhardt (Microsoft)
 - Robinson (MIT-LL)
 - Shah (USCB)





Constructing Graph Representations of Raw Data Source




- Raw data sources can contain information about multiple types of relations between entities
- The process of constructing a graph representation is specific to both the data source and the relationships represented by the graph

- The development time of parsing and graph construction algorithms can overwhelm the runtime of the algorithm



Graph Construction Using D4M



- D4M provides needed flexibility in the construction of large-scale, dynamic graphs at different resolutions and scopes



Graph Construction Using D4M: Parsing Raw Data Into Dense Tables



Proxy Logs

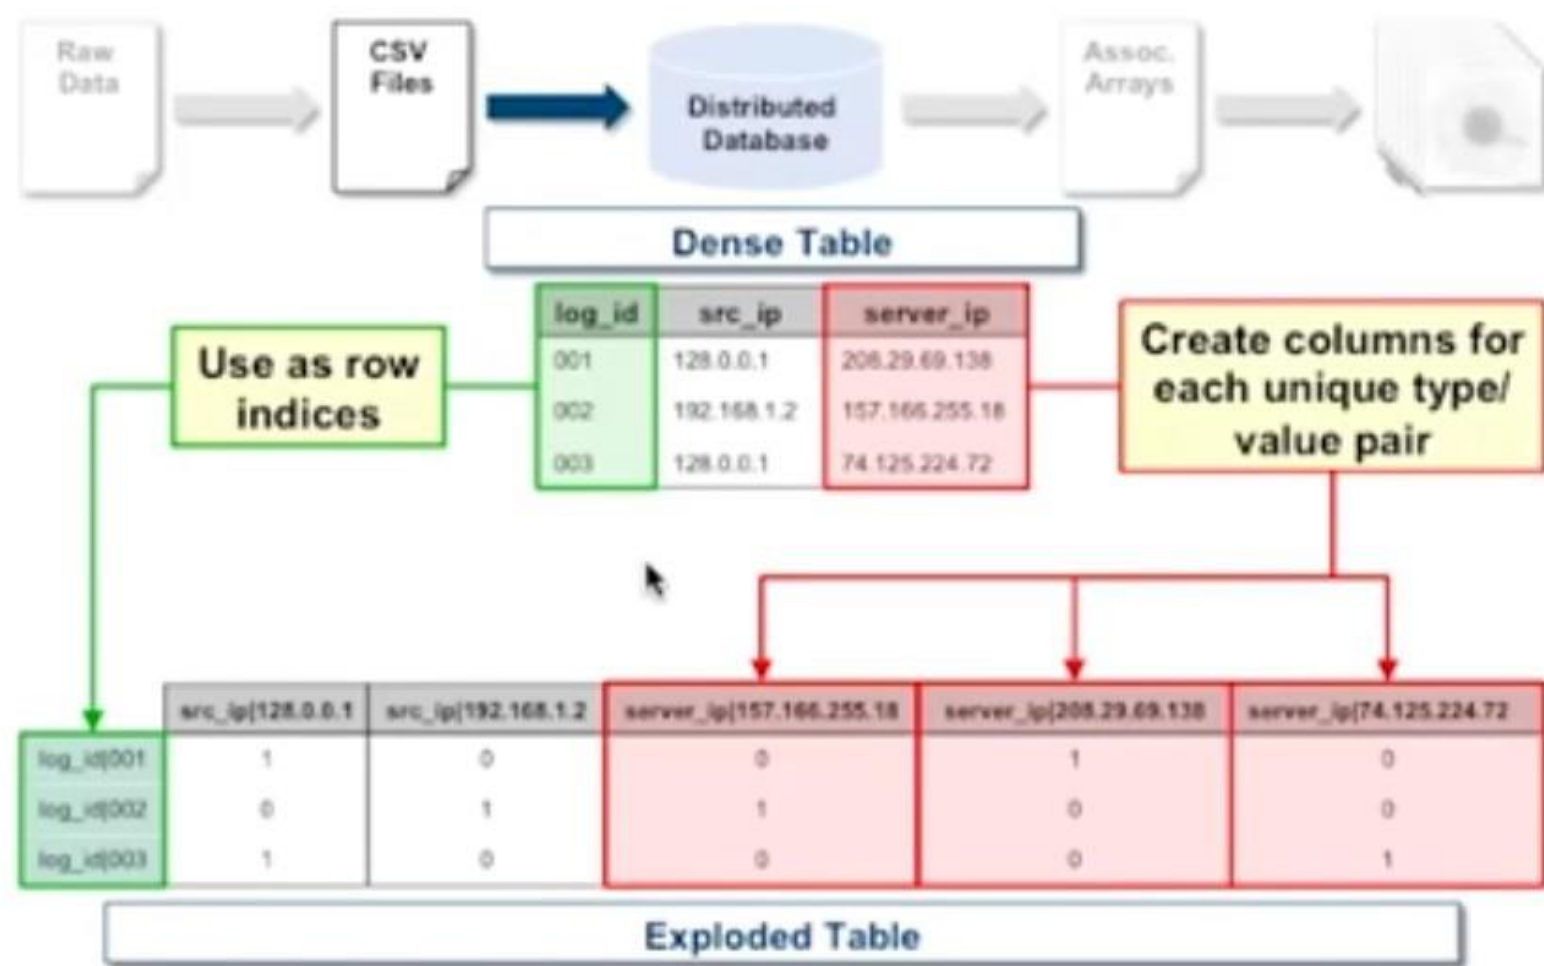
```
128.0.0.1 208.29.69.138 "-" [10/May/2011:09:52:53] "GET http://www.thedailybeast.com/ HTTP/1.1" 200  
1024 6192 "http://www.theatlantic.com/" "Mozilla/5.0 (X11; U; Linux x86_64; en-US; rv:1.9.2.13)  
Gecko/20101209 CentOS/3.6-2.el5.centos Firefox/3.6.13" "b1" - "text/html" "MITLAB" 0.523 "-"  
Neutral TCP_MISS  
192.168.1.1 157.166.255.18 "-" [12/May/2011:13:24:11] "GET http://www.cnn.com/ HTTP/1.1" 335 256  
10296 "-" "Mozilla/5.0 (X11; U; Linux x86_64; en-US; rv:1.9.2.13) Gecko/20101209 CentOS/  
3.6-2.el5.centos Firefox/3.6.13" "bu" - "text/html" "MITLAB" 0.784 "-" Neutral TCP_MISS  
...
```

Dense Table

| log_id | src_ip | server_ip | time_stamp | req_line | ... |
|--------|-------------|----------------|----------------------|--|-----|
| 001 | 128.0.0.1 | 208.29.69.138 | 10/May/2011:09:52:53 | GET http://www.thedailybeast.com/ HTTP/1.1 | ... |
| 002 | 192.168.1.2 | 157.166.255.18 | 12/May/2011:13:24:11 | GET http://www.cnn.com/ HTTP/1.1 | ... |
| 003 | 128.0.0.1 | 74.125.224.72 | 13/May/2011:11:05:12 | GET http://www.google.com/ HTTP/1.1 | ... |
| ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ |



Graph Construction Using D4M: Explode Schema





Graph Construction Using D4M: Storing Exploded Data as Triples



Exploded Table

| | src_ip[128.0.0.1] | src_ip[192.168.1.2] | server_ip[157.166.255.18] | server_ip[208.29.69.138] | server_ip[74.125.224.72] |
|-------------|-------------------|---------------------|---------------------------|--------------------------|--------------------------|
| log_id[001] | 1 | 0 | 0 | 1 | 0 |
| log_id[002] | 0 | 1 | 1 | 0 | 0 |
| log_id[003] | 1 | 0 | 0 | 0 | 1 |

D4M stores the triple data representing both
the exploded table and its transpose

Table Triples


| Row | Column | Value |
|-------------|---------------------------|-------|
| log_id[001] | src_ip[128.0.0.1] | 1 |
| log_id[001] | server_ip[208.29.69.138] | 1 |
| log_id[002] | src_ip[192.168.1.2] | 1 |
| log_id[002] | server_ip[157.166.255.18] | 1 |
| log_id[003] | src_ip[128.0.0.1] | 1 |
| log_id[003] | server_ip[74.125.224.72] | 1 |

Table Transpose Triples

| Row | Column | Value |
|---------------------------|-------------|-------|
| server_ip[157.166.255.18] | log_id[002] | 1 |
| server_ip[208.29.69.138] | log_id[001] | 1 |
| server_ip[74.125.224.72] | log_id[003] | 1 |
| src_ip[128.0.0.1] | log_id[001] | 1 |
| src_ip[128.0.0.1] | log_id[003] | 1 |
| src_ip[192.168.1.2] | log_id[002] | 1 |

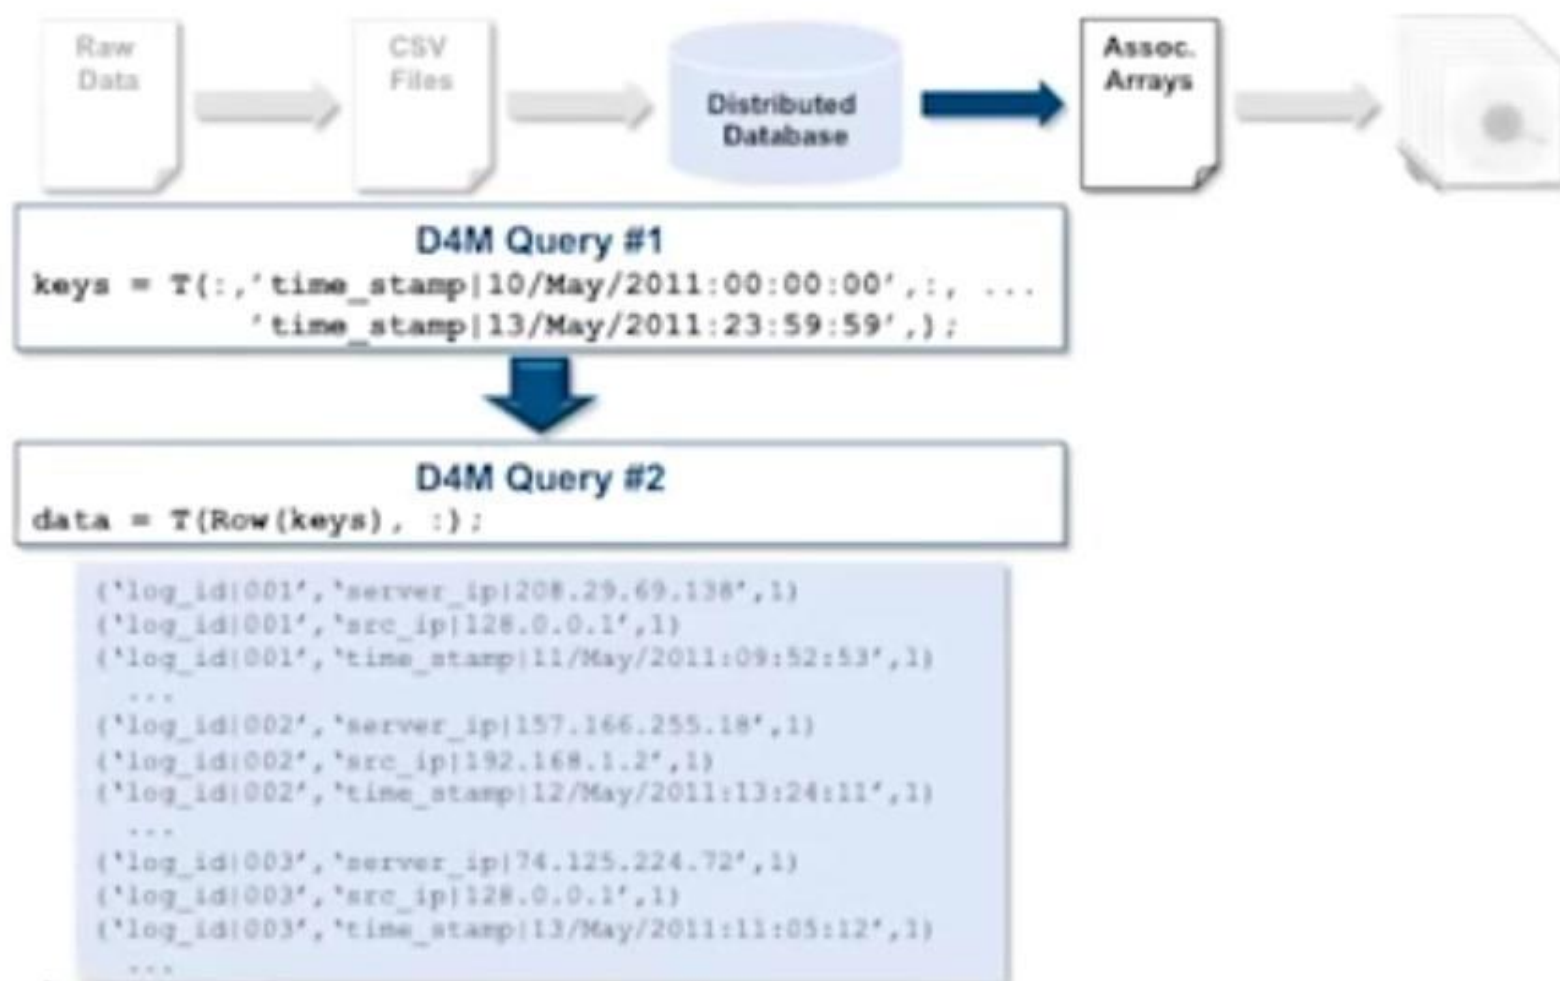


Graph Construction Using D4M: Construct Associative Arrays



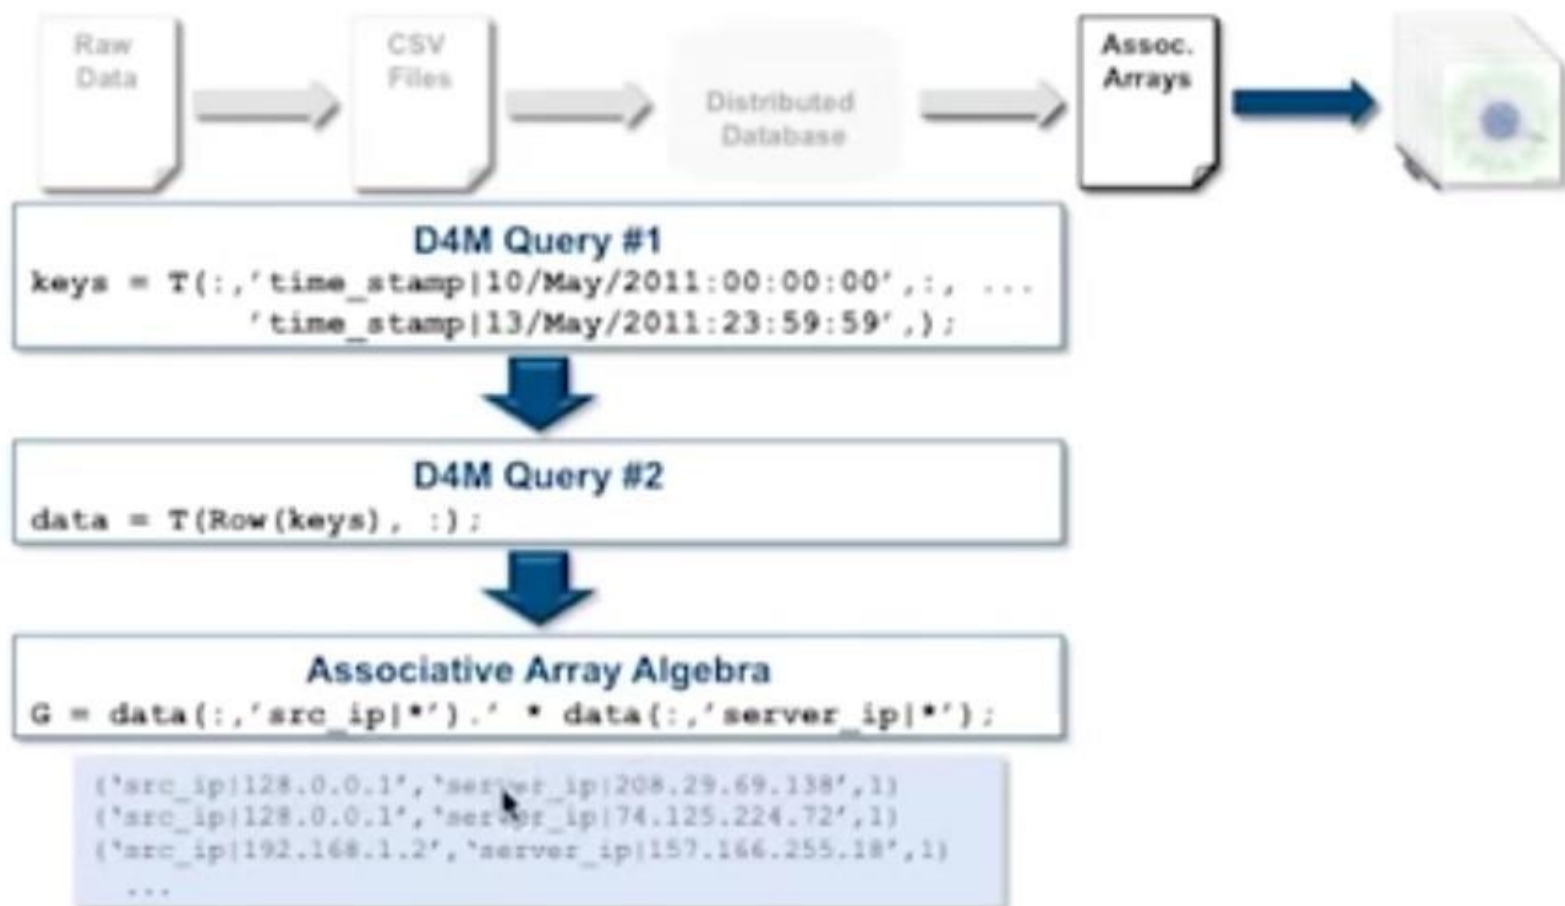


Graph Construction Using D4M: Construct Associative Arrays






Graph Construction Using D4M: Construct Associative Arrays





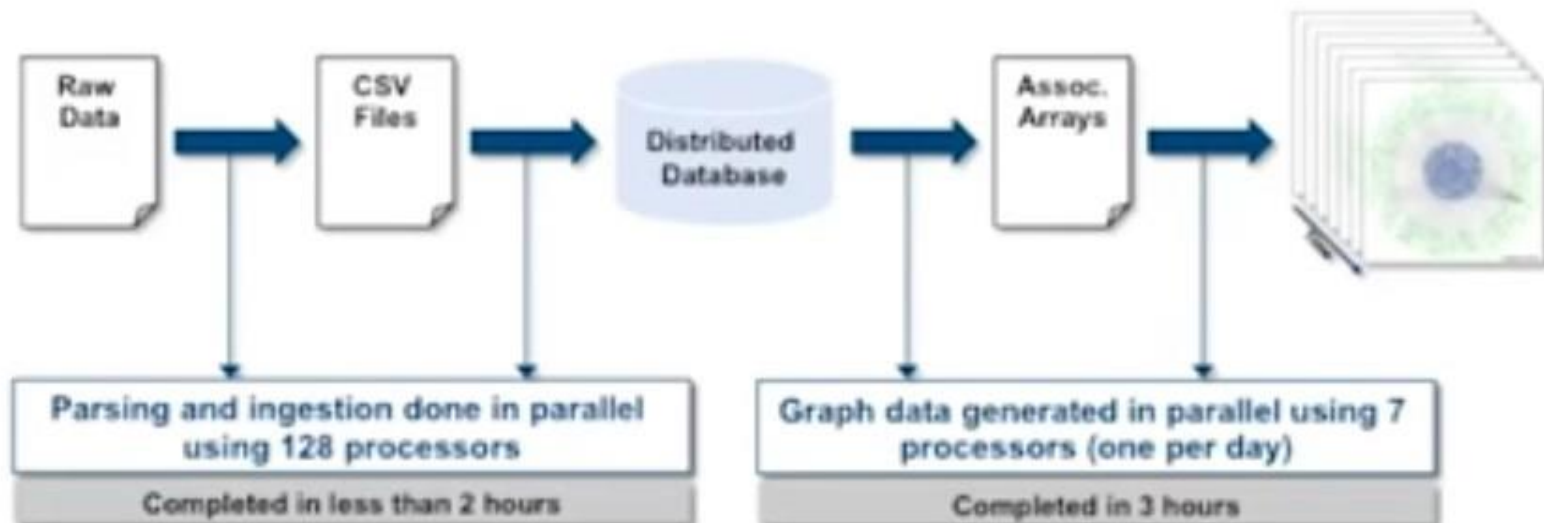
Graph Construction Using D4M: Construct Associative Arrays



- Graphs can be constructed with minimal effort using D4M queries and associative array algebra



Constructing Graph Representation of One Week's Worth of Proxy Data



- Ingested ~130 million proxy log records resulting in ~4.5 billion triples
- Constructed 604,800 secondwise source IP to server IP graphs
- Constructing graphs with different vertex types could be done without re-parsing or re-ingesting data

- Utilizing D4M could allow analysis to be run in nearly real-time (dependent on raw data availability)



Summary

- Big data is found across a wide range of areas
 - Document analysis
 - Computer network analysis
 - DNA Sequencing
- Currently there is a gap in big data analysis tools for algorithm developers
- D4M fills this gap by providing algorithm developers composable associative arrays that admit linear algebraic manipulation